

The production deployment of IPv6 on WLCG

J Bernier¹, S Campana², K Chadwick³, J Chudoba⁴, A Dewhurst⁵, M Eliáš⁴, S Fayer⁶, T Finnern⁷, C Grigoras², T Hartmann⁸, B Hoeft⁸, T Idiculla⁵, D P Kelsey⁵, F López Muñoz⁹, E Macmahon¹⁰, E Martelli², A P Millar⁷, R Nandakumar⁵, K Ohrenberg⁷, F Prelz¹¹, D Rand⁶, A Sciabà², U Tigerstedt¹², R Voicu¹³, C J Walker¹⁴ and T Wildish¹⁵

¹ IN2P3 Computing Center, 21 Avenue Pierre de Coubertin, F-69627 Villeurbanne Cedex, France

² CERN, CH-1211 Genève 23, Switzerland

³ Fermi National Accelerator Laboratory, Batavia, IL 60510, U.S.A.

⁴ Institute of Physics, Academy of Sciences of the Czech Republic Na Slovance 2 182 21 Prague 8, Czech Republic

⁵ STFC Rutherford Appleton Laboratory, Harwell Oxford, Didcot, Oxfordshire OX11 0QX, United Kingdom

⁶ Imperial College London, South Kensington Campus, London SW7 2AZ, United Kingdom

⁷ Deutsches Elektronen-Synchrotron, Notkestraße 85, D-22607 Hamburg, Germany

⁸ Karlsruher Institut für Technologie, Hermann-von-Helmholtz-Platz 1, D-76344 Eggenstein-Leopoldshafen, Germany

⁹ Port d'Informació Científica (PIC), Universitat Autònoma de Barcelona, Bellaterra (Barcelona), Spain. Also at Centro de Investigaciones Energéticas, Medioambientales y Tecnológicas, CIEMAT, Madrid, Spain

¹⁰ The University of Oxford, The Department of Physics, Denys Wilkinson Building, Keble Road, Oxford OX1 3RH, United Kingdom

¹¹ INFN, Sezione di Milano, via G. Celoria 16, I-20133 Milano, Italy

¹² CSC Tieteen Tietotekniikan Keskus Oy, P.O. Box 405, FI-02101 Espoo, Finland

¹³ California Institute of Technology, Pasadena, Ca 91125, U.S.A.

¹⁴ Queen Mary University of London, Mile End Road, London E1 4NS, United Kingdom

¹⁵ Princeton University, Jadwin Hall, Princeton, NJ 08544, U.S.A.

E-mail: david.kelsey@stfc.ac.uk, ipv6@hepox.org

Abstract. The world is rapidly running out of IPv4 addresses; the number of IPv6 end systems connected to the internet is increasing; WLCG and the LHC experiments may soon have access to worker nodes and/or virtual machines (VMs) possessing only an IPv6 routable address. The HEPiX IPv6 Working Group has been investigating, testing and planning for dual-stack services on WLCG for several years. Following feedback from our working group, many of the storage technologies in use on WLCG have recently been made IPv6-capable. This paper presents the IPv6 requirements, tests and plans of the LHC experiments together with the tests performed on the group's IPv6 test-bed. This is primarily aimed at IPv6-only worker nodes or VMs accessing several different implementations of a global dual-stack federated storage service. Finally the plans for deployment of production dual-stack WLCG services are presented.

1. Introduction

The world's Regional Internet Registries are rapidly running out of available IPv4 addresses and the general slow transition to IPv6 continues. The Worldwide Large Hadron Collider

BGP updates of IPv6 routes in 12 hours

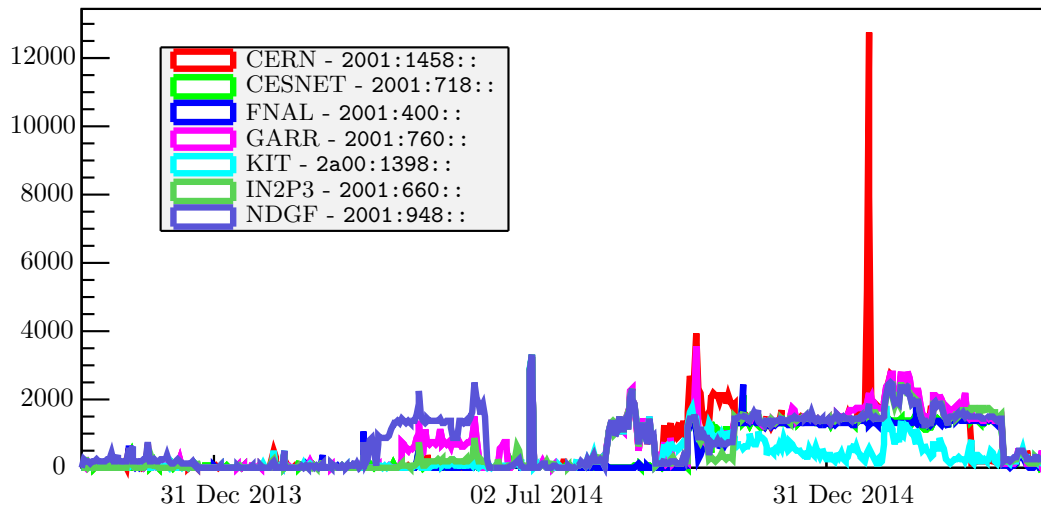


Figure 1. Number of routing topology updates involving the reference LIR of the working group testbed participants recorded by RIPEstat [4] every 12 hours since CHEP2013.

Grid (WLCG) and the LHC experiments may soon have access to worker nodes or virtual machines possessing only an IPv6-routable address. The HEPiX IPv6 Working Group [1] has been investigating the many issues feeding into the move to the use of IPv6 in HEP and WLCG. The group’s paper at CHEP2013 [2] described the aims of the group and the testing of dual-stack IPv6/IPv4 services that had been completed at that point. In the last 18 months the group has worked more closely with the 4 major LHC experiments and identified the main use case for the support of IPv6-only clients on WLCG. The groups activities, including testing of dual-stack data storage services, during the last 18 months are presented in this paper together with its future plans.

2. Status and hurdles of the worldwide IPv4→IPv6 transition

We now offer a quick panorama of the IPv6 statistical trends since the last CHEP conference and identify two factors that may be currently limiting the IPv6 adoption rate.

2.1. Survey of available statistical data

An analysis of the available statistical data collected since 2013 by the Regional Internet Registries [5–9] and by major internet service providers [10, 11] shows a steady, but still polynomial growth in the volume of IPv6 traffic from roughly 2% up to 6% of the total. Not everything has been progressing at a relaxed pace, though: we collected from RIPEstat [4] into Figure 1 the rate of BGP routing topology updates that affects the IPv6 /32 prefixes that serve our working group testbed. We notice a significantly increased activity rate over the last year. We can probably read this as the effect of organising and troubleshooting efficient, production-proof routing for our community.

On the IPv4 address exhaustion side, the actual availability of assignable IPv4 addresses has remained substantially stable around the 18 million mark at RIPE (Europe). AFRINIC (Africa) still has all the 16 million addresses in the last assigned /8 network available, APNIC (Asia-Pacific) has 12 million address left, ARIN (North America) has 4 million and LACNIC (South America) is in the most critical state with 3 million IPv4 addresses left. All Regional Internet

Registries are now implementing IPv4 request limits and ‘soft landing’ allocation schemes.

2.2. Slow IPv6 adoption progress rate: why?

While it is in the best interest of a successful transition to IPv6 that no show-stoppers be found and the process keeps moving *forward*, one may wonder about the causes that prevent a more rapid adoption pace.

Transport and provider issues have to be excluded right away: most local registries, including all of the national research networks, have been providing IPv6 transport for 7 years or more [3]. Performance issues also have to be ruled out: available studies, especially the ones carried out during the “world IPv6 launch” in 2012 [18] show performance on the two stacks to be comparable within statistics. The reality of the IPv4 address depletion should also be widely perceived by now, as many regional registries are handling the *final* IPv4 assignment to local registries.

Two residual classes of factors may be quenching IPv6 adoption, one affecting network administrators and one affecting application developers:

- (i) The IPv4/v6 difference in address allocation schemes, the pivot role of ICMPv6 Router Advertisements, the short-term need to implement measures to counter rogue Router Advertisements (see RFC6104, [17]) are all adding to the already sizeable initial investment of implementing monitoring and security tools for IPv6.

Also, existing IPv6 code in the Operating Systems and related tools often shows by inspection not to have undergone full coverage testing: a phase of initial fault finding and patching is foreseen and feared.

- (ii) Apart from the syntactical differences in IPv6 addresses (e.g. parsing ‘defa’ in *default* as a hex digit), and the need to label, sort and pick IPv4 vs IPv6 addresses, a large *semantical* change is needed in applications supporting IPv6. Every network endpoint on the public IPv6 network has *at least* two IPv6 addresses assigned (global and link-local), and possibly more. Applications have therefore to *always* deal with multi-homed network endpoints, which means a complex $1 \rightarrow n$ change for many of them. The status of porting to IPv6 of many applications of interest for our community is good [16], but the related development effort cannot be underestimated.

3. The survey of IPv6 readiness at WLCG sites

In the summer of 2014 a survey of all WLCG sites was performed by the IPv6 working group. This asked a few simple questions to determine the site readiness for IPv6 and also whether they foresaw running out of IPv4 address space in the coming years.

Table 1. Site IPv6-readiness

Type of Site	Answered	IPv6 now	IPv6 soon	No IPv6 plans	Lack of IPv4
Tier 0/1	14	8	4	2	2
Tier 2	100	24	14	62	10

Approximately two-thirds of the WLCG sites responded and the broad conclusions of the survey are summarised in table 1. “IPv6 now” means that the site had IPv6 connectivity at the time of the survey. “IPv6 soon” means that such connectivity is planned to be available within two years. “No IPv6 plans” means that either the site has not started planning or the

planned date is more than 2 years away. "Lack of IPv4" means that the site has already run out of IPv4-routable addresses or foresees this to happen within the next 2 years.

The main conclusions of the survey are that most Tier 1 sites are or will soon be ready, whereas approximately 60% of the Tier 2 sites have not yet started their planning for IPv6. Moreover the fact that about 10% of the sites foresee problems with the imminent lack of routable IPv4 addresses means that WLCG must consider moving to the use of dual-stack IPv6/IPv4 services as quickly as possible.

4. LHC Experiment requirements and main use case

The shortage of available IPv4 addresses implies that there is a significant possibility that new large computing facilities will not be able to give IPv4 addresses to all of the machines in their network. The most likely consequence is that for these sites, worker nodes – which constitute the largest fraction of independent computing nodes will have purely IPv6 network addresses. Hence, the main use case for the LHC experiments is to enable jobs to run on these machines, access their software areas and input data and upload their outputs to various grid storages or services as needed.

The LHC experiments generally assume [12] that the storage on different sites and supporting middleware [13] like the LFC will either be directly dual-stack, or support dual stack operation in some way, enabling seamless access to the storage as needed for either downloading or saving. For example, it is expected that dual-stack squid proxies will be needed for CVMFS and xrootd will soon be dual-stack, to handle storage technologies like Castor which will not be IPv4 only. The servers that the LHC experiments use to handle the grid infrastructure are / will also be dual-stack.

As an example of the above, we look at LHCb [14] which is the experiment on the LHC, optimised for studying beauty and charm physics. LHCb uses the DIRAC [?] interware to manage its grid operations. The DIRAC software was coded to be able to handle both IPv4 and IPv6 addresses in late 2014, with the modifications being easy enough to make by non-expert programmers. While testing the processes on a dual-stack machine, it was found that there was a significant number of connections which were not going through to the servers which was finally traced back to a missing `enable_ipv6` option to compile python by an external provider causing errors in identifying IPv6 addresses. Using a new version of the library, the problem with dropped connections went away and testing DIRAC will restart soon.

In general, testing of the grid middleware by the different experiments is going ahead as fast as possible given the manpower and time constraints of the LHC startup and the immediate issues with handling purely IPv6 worker nodes are expected to be sorted by sometime in 2016.

5. Testbed operation: testing FTS3/dCache

5.1. The Transfer Testbed

The transfer testbed was upgraded in March 2015. Until then, it operated with gridFTP transfers between all sites, providing a low-level test of connectivity and functionality for the almost two years that it ran. Since March 2015 the testbed uses FTS3 [19] to initiate the transfers, moving up the middleware stack. Since FTS3 is used for the vast majority of experiment transfers in WLCG this provides an important full-stack test.

At the present time, the testbed consists of 7 storage elements at sites distributed around Europe. One is IPv6-only, the rest are all dual-stack. All the SEs are running dCache. Most are stable installations, but one (DESY) is rebuilt every morning with the latest patches from dCache, providing a valuable regression-test for both the dCache and IPv6 teams.

As before, each site serves as both a source and a destination, with each source sending a 1 GB file to each destination. The file-size is validated at the destination using **gfal-ls**, then the destination is cleaned with **gfal-rm** and the transfer duration is recorded. Then the cycle

is repeated after a short delay, to avoid abusing the hardware/network with too much traffic. Physical file names are specified using the SRM protocol.

Two FTS3 servers are deployed for the testbed, one at Imperial College and one at KIT, though currently only the one at KIT is used.

Figure. 2 shows the transfers in the FTS3 testbed so far. Most sites transfer efficiently in both directions, but the effect of the firewall at KIT on inbound traffic can be clearly seen.

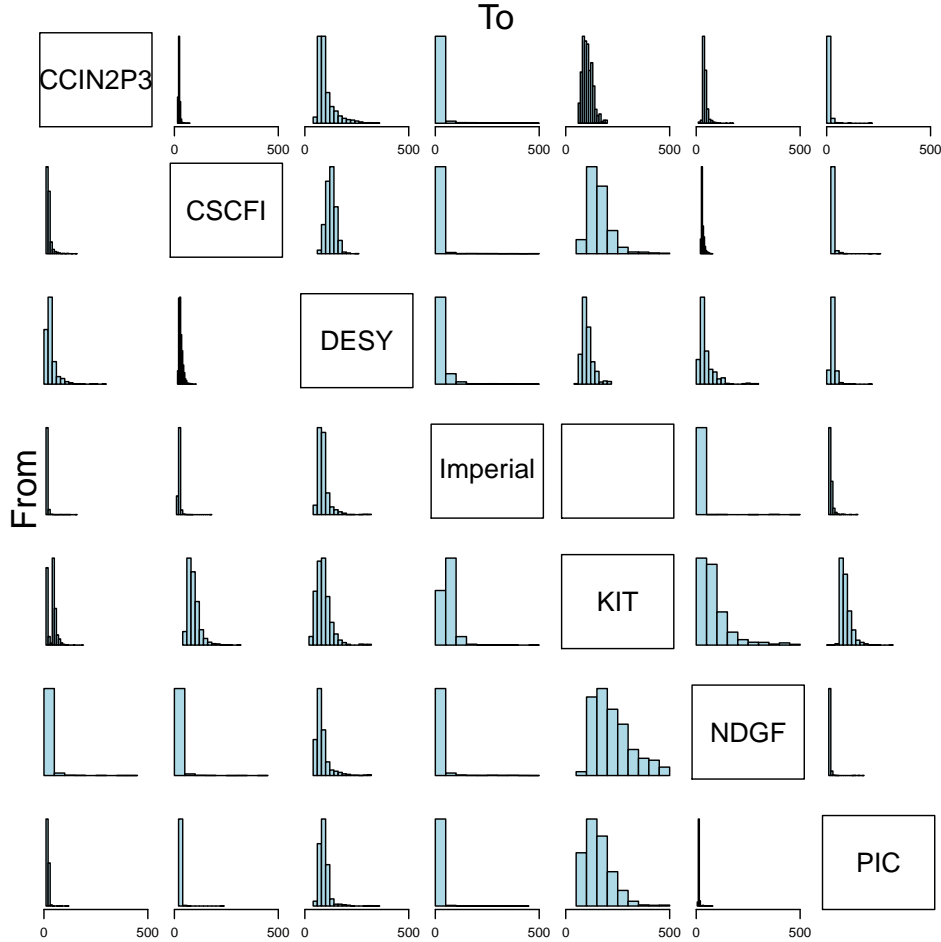


Figure 2. The FTS3 transfer testbed. Rows show transfers from the named site, columns show transfers to the destination. The horizontal axis for all plots is fixed at 500 seconds, i.e. transfers that proceed at less than 2 MB/sec will overflow.

5.2. FTS3 server and dCache SE at KIT

For managing file transfers between sites, an FTS3 instance is setup at KIT. Furthermore, a storage element based on dCache 2.10 on Scientific Linux is created. Both instances are rolled out on physical machines.

5.2.1. FTS3 FTS3 supports IPv6 in its baseline version. The service has to be bound to IPv6 locally in the FTS3 conguration (IP=::) and to be enabled explicitly for gfal2 (IPV6=true). The host is available in the DNS as default dual-host, with IP4 and IPv6 announced in the A and

AAAA records, and also with IPv4-only or IPv6-only names with an -ipv4 or -ipv6 appendix, respectively. Thus, all aliases have to be included in the host certificate.

File transfers are successfully brokered by the FTS3 instance via IPv4 and IPv6 between the sites. Since most FTS3 instances in production use a separated database instance for performance and failsafe reasons, moving the database to a dedicated machine was tested as well. For the SQL db backends supported by FTS3, IPv6 support had been implemented in MySQL v5.5.3 and MariaDB v5.5.35, which are not available in the baseline SL6 repositories. MariaDB was installed on a dedicated host with version 5.5.42. After binding mysql locally to IPv6 as well ([mysqld] bind-address = ::), the database could be connected remotely with an IPv6-ready mysql client. For the FTS3 service to connect to the remote database via IPv6, the address has to be escaped explicitly, i.e., encapsulating the IP as [IP6]:PORT/fts3 and may depend on the version of the database library used by the FTS3 service.

5.2.2. dCache based SE A dCache instance is setup on a dedicated host. The main hurdle for file transfer access was a reverse lookup by the instance when receiving a file request. After explicitly setting the dual-stack, IPv4-only, and IPv6-only host names in the dCache configuration (srm.net.local-hosts=hostname-ipv4,ipv6) and the general hosts file, the storage element is accessible via IPv4 and IPv6 as well.

6. IPv6 readiness of storage technology for WLCG

6.1. Storage technology

The old testbed was based on GridFTP [20] [21] transfers between small storage elements. The software used on the servers were Globus gridftp, DPM (the Disk Pool Manager) and dCache. All of these needed configuration changes from the default to work with IPv6.

- DPM needed a new MySQL (v5.5+) and several configuration changes to do any transfers over IPv6.
- dCache got GridFTP working in version 2.9.4, but needed no configuration changes.
- StoRM needs configuration changes but works if it's new enough.
- XRootD got first IPv6 support in version 4.0.0, it is fully dual-stack since. Version 4.1.0 brought a lot of fixes for the IPv6 functionality.
- Globus gridftp needs configuration changes.

For the FTS3-based testbed the SRM protocol [22] was selected for initiating transfers as it is more real-life production-like than the old gridftp-based one.

All sites participating in the new testbed selected dCache, as it had matured into the only full-stack storage system to fully support dualstack and IPv6-only setups. Even with full support for IPv6 as default, the testbeds setup with multiple hostnames (dualstack, IPv4-only and IPv6-only) needed configuration changes for the SRM door for it to recognize that it had multiple hostnames.

A new addition was also a IPv6-only storage element, to test how well clients and FTS3 could handle a storage element that is not dual-stacked. This required very little configuration but would not work without the latest release of dCache, 2.12.

7. IPv6 perfSONAR measurements

The WLCG has adopted the perfSONAR toolkit [23] for the monitoring of its network infrastructure and this project is being coordinated by the WLCG Network and Transfer Metrics group [24]. The WLCG perfSONAR configuration system operates around groups of sites with a common purpose and these are known as meshes. For example, there are meshes for each WLCG country group (e.g. UK, DE, FR etc.) or experiment such as USATLAS or USCMS or

network groupings such as LHCONE and LHCOPN. Until recently testing between members of the groups was configured using JSON files held on a web-server at CERN specifying the group members and test parameters. A site administrator configuring a perfSONAR host for a mesh needed to add the URL of the JSON file corresponding to that mesh into a configuration file on the perfSONAR host. This worked but required effort from all site administrators involved. The mesh configuration system has recently undergone development. This is described in detail in [24], but briefly, the system has evolved from the use of the manually configured JSON files to a more automated system which is significantly easier to configure. Each perfSONAR host now has a so-called auto-mesh URL e.g.

<https://myosg.grid.iu.edu/pfmesh/mine/hostname/psum01.aglt2.org>

containing configuration details for the meshes that the perfSONAR host has been added to. The meshes are maintained using a mesh-configuration GUI provided by the US Open Sciences Grid (OSG) using data collected from both the GOCDB and OSG Information Management System (OIM). The perfSONAR toolkit has the ability to monitor both IPv4 and IPv6 network connectivity. Consequently, in addition to the meshes mentioned above, we have added a mesh containing perfSONAR hosts known to have both IPv4 and IPv6 connectivity, i.e. a dual-stack mesh. The mesh tests throughput and latency between hosts over both IPv4 and IPv6. Results are available from the web sites of the relevant perfSONAR hosts, for example the perfSONAR bandwidth host at WLCG site UKI-SOUTHGRID-OX-HEP at the University of Oxford:

<http://t2ps-bandwidth.physics.ox.ac.uk/toolkit/>

8. Status of LHCOPNv6/LHCONEv6

In November and December 2014 the Grid Deployment Board requested tier-1s to join the HEPiX-IPv6 working group and further encouraged sites to move their production endpoints to dual stack even if this required concessions of the quotation of their site reliability and site availability. The proposal of the LHC experiment Atlas was to

- request that all Tier-1s provide, besides an IPv6 peering to LHCOPN, a dual stack PerfSONAR machine by first of April 2015
- request that Tier-2s provide, besides an IPv6 peering to their LHCONE connection, a dual stack PerfSONAR machine by August 2015.

At the last LHC[OPN/ONE] meeting in February this year a proposal containing the IPv6 request was put forward. There were no objections from the site representatives nor from the NRENs to this proposal. The following Tier-1 sites are actively announcing an IPv6 peering to LHCOPN: CH-CERN, DE-KIT, ES-PIC, FR-CCIN2P3, NDGF, NL-T1. IT-INFN-CNAF is in the process of preparing the IPv6 peering. The group of IPv6 peers over LHCONE is currently even smaller: besides CH-CERN this are the two sites CEA SACLAY and FR-CCIN2P3. The ipv6 peerings are reflected at the PerfSONAR dualstack dashboard url:

<http://maddash.aglt2.org/maddash-webui/index.cgi?dashboard=Dual-Stack%20Mesh%20Config>

. It implies that there are still some LHC tier-1 sites, more than two month behind the deadline agreed upon, not announcing their ipv6 cidr to LHCONE.

9. Outlook and future plans

The HEPiX IPv6 working group has made good progress during the last 18 months. It has been demonstrated that access to remote dual-stack federated data storage services in a production-like environment functions well with FTS, SRM and dCache. During the remainder of 2015 tests on other storage technologies will be performed. Members of the group will deploy dual-stack services on more production instances of storage services and other essential central services to enable the proper measurement of data transfer performance over IPv6 and to demonstrate to

the experiments and the WLCG management that it is safe to migrate to dual-stack. The deployment of IPv6 peering on LHCOPN/LHCONE and dual-stack perfSONAR instances will be tracked and encouraged. Both of these are pre-requisites to the wider deployment of production dual-stack services. In parallel with these activities the group also aims to provide more training sessions and guidance on best practice in the management of IPv6 services and site operations. Once there are a sufficient number of dual-stack services deployed on WLCG it will be possible to support the use of IPv6-only clients within the production infrastructure.

References

- [1] <http://hepex-ipv6.web.cern.ch>
- [2] S Campana et al 2014 J. Phys.: Conf. Ser. **513** 062026
- [3] The status of IPv6 readiness of the local Internet Registries that refer to RIPE is monitored in the “IPv6 RIPEness” pages (<https://ipv6ripeness.ripe.net/> at the time of writing).
- [4] RIPE provides access to a comprehensive set of statistical data via <https://stat.ripe.net/data>. This data set is also used for the online diagnostics and statistical tools (e.g. BGPlay).
- [5] Specific IPv6 trend graphs for the RIPE LIRs can be accessed at <https://labs.ripe.net/statistics/?tags=ipv6>.
- [6] A collection of trend plots collected by ARIN can be found here: <https://www.arin.net/knowledge/statistics/>.
- [7] A somewhat out-of-date set of IPv6 statistics data gathered by APNIC can be found here: <https://labs.apnic.net/ipv6-measurement/>.
- [8] The LACNIC IPv6 portal is at: <http://portalipv6.lacnic.net/en/>.
- [9] A collection of AFRINIC IPv6 resources can be found here: <http://www.afrinic.net/en/services/statistics/ipv6-resources>.
- [10] The historical statistics of IPv6 traffic on Akamai usare are recorded at <http://www.akamai.com/ipv6/> at the time of writing.
- [11] The historical statistics of IPv6 client traffic on Google are recorded at <http://www.google.com/intl/en/ipv6/statistics.html> at the time of writing.
- [12] WLCG pre-GDB - IPv6 Workshop (A. Dewhurst et al.), <https://indico.cern.ch/event/313194/session/1/contribution/3>
- [13] <https://wiki.egi.eu/wiki/Middleware.products.verified.for.the.support.of.IPv6> , <http://hepex-ipv6.web.cern.ch/wlcg-applications>
- [14] LHCb Collab. (A. Alves et al.), JINST 3, S08005 (2008), <http://dx.doi.org/10.1088/1748-0221/3/08/S08005>
- [15] DIRAC: a community grid solution (A. Tsaregorodtsev et al.), 2008 J. Phys.: Conf. Ser. 119 062048, <http://iopscience.iop.org/1742-6596/119/6/062048>
- [16] Our working group tracks the IPv6 readiness of applications of interest to WLCG at <http://hepex-ipv6.web.cern.ch/wlcg-applications>.
- [17] All Internet Engineering Task Force Requests For Comments (RFC) documents are available from URLs such as <http://www.ietf.org/rfc/rfcNNNN.txt> where NNNN is the RFC number, for example <http://www.ietf.org/rfc/rfc2460.txt>
- [18] Plonka D, Barford P, Assessing performance of Internet services on IPv6, 2013 IEEE Symposium on Computers and Communications (ISCC), doi:10.1109/ISCC.2013.6755050
- [19] A Ayllon and M Salichos and M K Simon and O Keeble *FTS3: New Data Movement Service For WLCG*, J. Phys.: Conf. Ser. 2014, **513** 3 032081, doi:10.1088/1742-6596/513/3/032081, <http://dx.doi.org/10.1088/1742-6596/513/3/032081>,
- [20] <https://www.ogf.org/documents/GFD.20.pdf>
- [21] <https://www.ogf.org/documents/GFD.21.pdf>
- [22] <http://sdm.lbl.gov/srm-wg/doc/SRM.v2.2.html>
- [23] B Tierney, J Metzger, J Boote, A Brown, M Zekauskas J Zurawski, M Swany, M Grigoriev, perfSONAR: Instantiating a Global Network Measurement Framework, 4th Workshop on Real Overlays and Distributed Systems (ROADS09) Co-located with the 22nd ACM Symposium on Operating Systems Principles (SOSP), January 1, 2009.
- [24] S McKee, M Babik et al. Integrating network and transfer metrics to optimize transfer efficiency and experiment workflows, CHEP2015, Journal of Physics: Conference Series.